



Working With RDDs in Spark

Chapter 11



Course Chapters

1	Introduction	Course Introduction
2	Introduction to Hadoop and the Hadoop Ecosystem	Introduction to Hadoop
3	Hadoop Architecture and HDFS	
4	Importing Relational Data with Apache Sqoop	Importing and Modeling Structured Data
5	Introduction to Impala and Hive	
6	Modeling and Managing Data with Impala and Hive	
7	Data Formats	
8	Data File Partitioning	
9	Capturing Data with Apache Flume	Ingesting Streaming Data
10	Spark Basics	Distributed Data Processing with Spark
11	Working with RDDs in Spark	
12	Aggregating Data with Pair RDDs	
13	Writing and Deploying Spark Applications	
14	Parallel Processing in Spark	
15	Spark RDD Persistence	
16	Common Patterns in Spark Data Processing	
17	Spark SQL and DataFrames	
18	Conclusion	Course Conclusion

Working With RDDs

In this chapter you will learn

- **How RDDs are created from files or data in memory**
- **How to handle file formats with multi-line records**
- **How to use some additional operations on RDDs**

Chapter Topics

Working With RDDs in Spark

Distributed Data Processing with Spark

- **Creating RDDs**
- Other General RDD Operations
- Conclusion
- Homework: Process Data Files with Spark

RDDs

- **RDDs can hold any type of element**
 - Primitive types: integers, characters, booleans, etc.
 - Sequence types: strings, lists, arrays, tuples, dicts, etc. (including nested data types)
 - Scala/Java Objects (if serializable)
 - Mixed types
- **Some types of RDDs have additional functionality**
 - Pair RDDs
 - RDDs consisting of Key-Value pairs
 - Double RDDs
 - RDDs consisting of numeric data

Creating RDDs From Collections

- You can create RDDs from collections instead of files
 - `sc.parallelize(collection)`

```
> myData = ["Alice", "Carlos", "Frank", "Barbara"]  
> myRdd = sc.parallelize(myData)  
> myRdd.take(2)  
['Alice', 'Carlos']
```

- Useful when
 - Testing
 - Generating data programmatically
 - Integrating

Creating RDDs from Files (1)

- **For file-based RDDs, use `SparkContext.textFile`**
 - Accepts a single file, a wildcard list of files, or a comma-separated list of files
 - Examples
 - `sc.textFile("myfile.txt")`
 - `sc.textFile("mydata/*.log")`
 - `sc.textFile("myfile1.txt,myfile2.txt")`
 - Each line in the file(s) is a separate record in the RDD
- **Files are referenced by absolute or relative URI**
 - Absolute URI:
 - `file:/home/training/myfile.txt`
 - `hdfs://localhost/loudacre/myfile.txt`
 - Relative URI (uses default file system): `myfile.txt`

Creating RDDs from Files (2)

- **textFile** maps each line in a file to a separate RDD element

```
I've never seen a purple cow.\nI never hope to see one;\nBut I can tell you, anyhow,\nI'd rather see than be one.\n
```



I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

- **textFile** only works with line-delimited text files
- What about other formats?

Input and Output Formats (1)

- **Spark uses Hadoop `InputFormat` and `OutputFormat` Java classes**
 - Some examples from core Hadoop
 - **`TextInputFormat` / `TextOutputFormat`** – newline delimited text files
 - **`SequenceInputFormat` / `SequenceOutputFormat`**
 - **`FixedLengthInputFormat`**
 - Many implementations available in additional libraries
 - e.g. **`AvroInputFormat` / `AvroOutputFormat`** in the Avro library

Input and Output Formats (2)

- Specify any input format using `sc.hadoopFile`
 - or `newAPIHadoopFile` for New API classes
- Specify any output format using `rdd.saveAsHadoopFile`
 - or `saveAsNewAPIHadoopFile` for New API classes
- `textFile` and `saveAsTextFile` are convenience functions
 - `textFile` just calls `hadoopFile` specifying `TextInputFormat`
 - `saveAsTextFile` calls `saveAsHadoopFile` specifying `TextOutputFormat`

Whole File-Based RDDs (1)


- **`sc.textFile`** maps each line in a file to a separate RDD element
 - What about files with a multi-line input format, e.g. XML or JSON?
- **`sc.wholeTextFiles(directory)`**
 - Maps entire contents of each file in a directory to a single RDD element
 - Works only for small files (element must fit in memory)

file1.json

```
{
  "firstName": "Fred",
  "lastName": "Flintstone",
  "userid": "123"
}
```

file2.json

```
{
  "firstName": "Barney",
  "lastName": "Rubble",
  "userid": "234"
}
```



<code>(file1.json, {"firstName": "Fred", "lastName": "Flintstone", "userid": "123"})</code>
<code>(file2.json, {"firstName": "Barney", "lastName": "Rubble", "userid": "234"})</code>
<code>(file3.xml, ...)</code>
<code>(file4.xml, ...)</code>

Whole File-Based RDDs (2)

```
> import json
> myrdd1 = sc.wholeTextFiles(mydir)
> myrdd2 = myrdd1
  .map(lambda (fname,s): json.loads(s))
> for record in myrdd2.take(2):
>     print record["firstName"]
```

```
> import scala.util.parsing.json.JSON
> val myrdd1 = sc.wholeTextFiles(mydir)
> val myrdd2 = myrdd1
  .map(pair => JSON.parseFull(pair._2).get.
    asInstanceOf[Map[String,String]])
> for (record <- myrdd2.take(2))
  println(record.getOrElse("firstName",null))
```

Output:

Fred
Barney

Chapter Topics

Working With RDDs in Spark

Distributed Data Processing with Spark

- Creating RDDs
- **Other General RDD Operations**
- Conclusion
- Homework: Process Data Files with Spark

Some Other General RDD Operations

■ Single-RDD Transformations

- **flatMap** – maps one element in the base RDD to multiple elements
- **distinct** – filter out duplicates
- **sortBy** – use provided function to sort

■ Multi-RDD Transformations

- **intersection** – create a new RDD with all elements in both original RDDs
- **union** – add all elements of two RDDs into a single new RDD
- **zip** – pair each element of the first RDD with the corresponding element of the second

Example: flatMap and distinct

Python

```
> sc.textFile(file) \
    .flatMap(lambda line: line.split()) \
    .distinct()
```

Scala

```
> sc.textFile(file).
    flatMap(line => line.split(' ')).
    distinct()
```

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

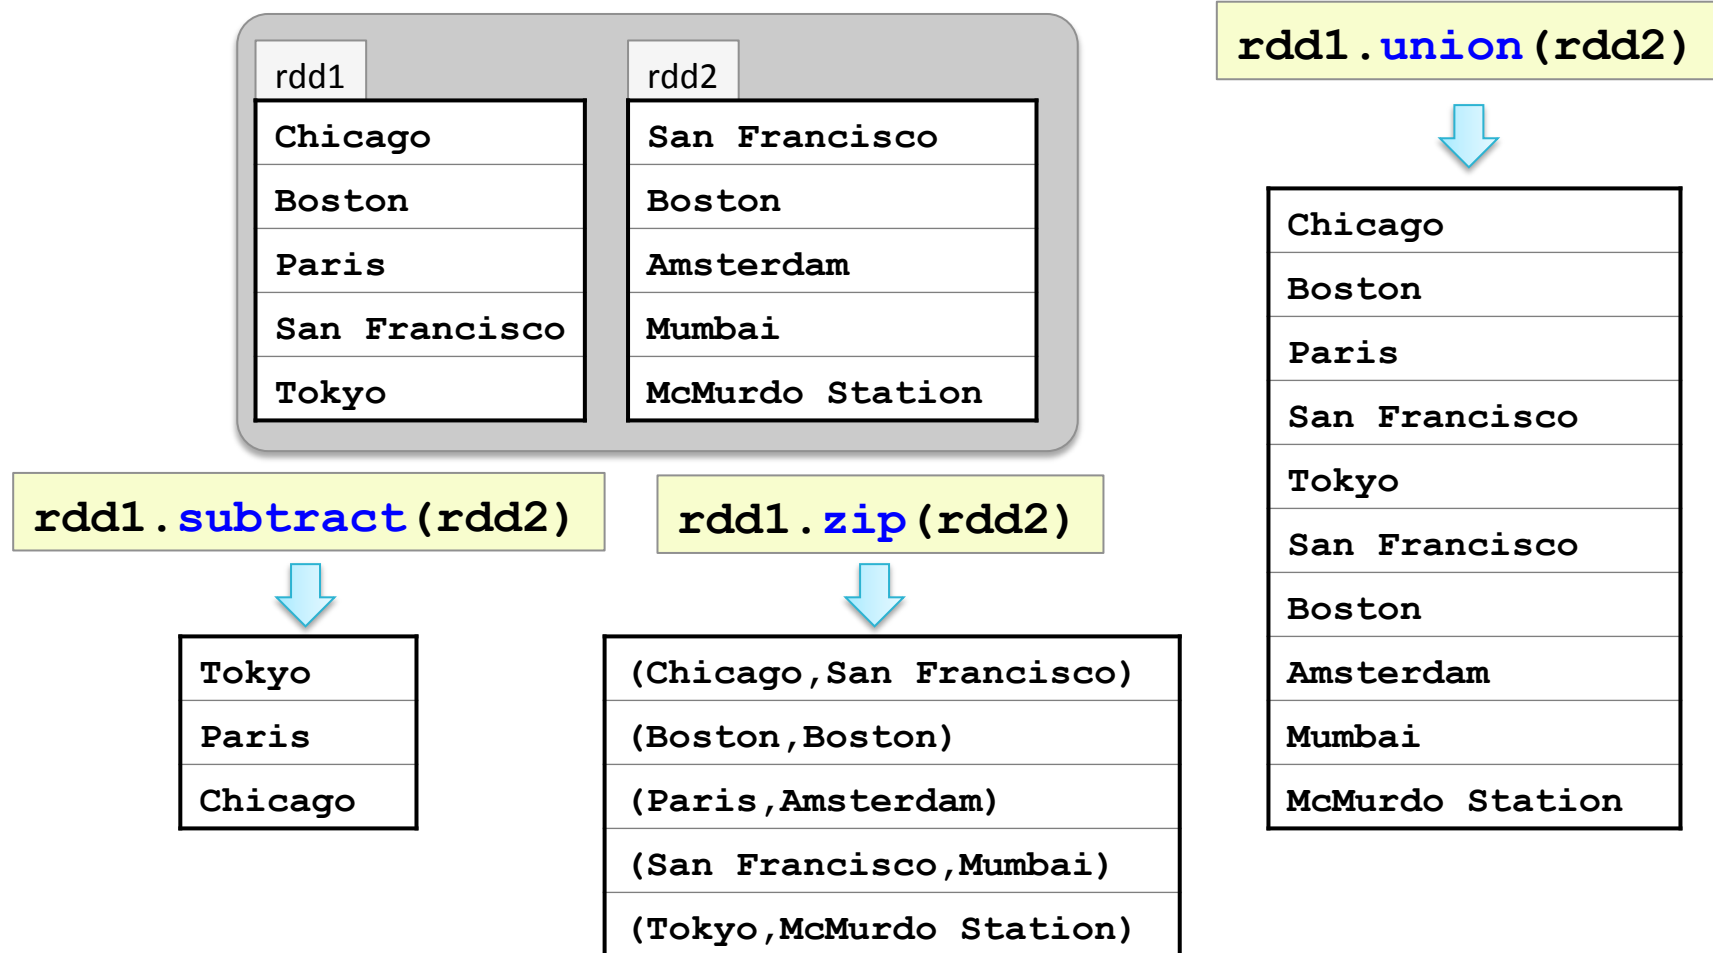


I've
never
seen
a
purple
cow
I
never
hope
to
...



I've
never
seen
a
purple
cow
I
hope
to
...

Examples: Multi-RDD Transformations



Some Other General RDD Operations

- **Other RDD operations**

- **first** – return the first element of the RDD
- **foreach** – apply a function to each element in an RDD
- **top (*n*)** – return the largest *n* elements using natural ordering

- **Sampling operations**

- **sample** – create a new RDD with a sampling of elements
- **takeSample** – return an array of sampled elements

- **Double RDD operations**

- Statistical functions, e.g., **mean**, **sum**, **variance**, **stdev**

Chapter Topics

Working With RDDs in Spark

Distributed Data Processing with Spark

- Creating RDDs
- Other General RDD Operations
- **Conclusion**
- Homework: Process Data Files with Spark

Essential Points

- **RDDs can be created from files, parallelized data in memory, or other RDDs**
- **`sc.textFile` reads newline delimited text, one line per RDD record**
- **`sc.wholeTextFile` reads entire files into single RDD records**
- **Generic RDDs can consist of any type of data**
- **Generic RDDs provide a wide range of transformation operations**

Chapter Topics

Working With RDDs in Spark

Distributed Data Processing with Spark

- Creating RDDs
- Other General RDD Operations
- Conclusion
- **Homework: Process Data Files with Spark**

Homework: Process Data Files with Spark

- **In this homework assignment you will**
 - Process a set of XML files using **wholeTextFiles**
 - Reformat a dataset to standardize format (bonus)
- **Please refer to the Homework description**